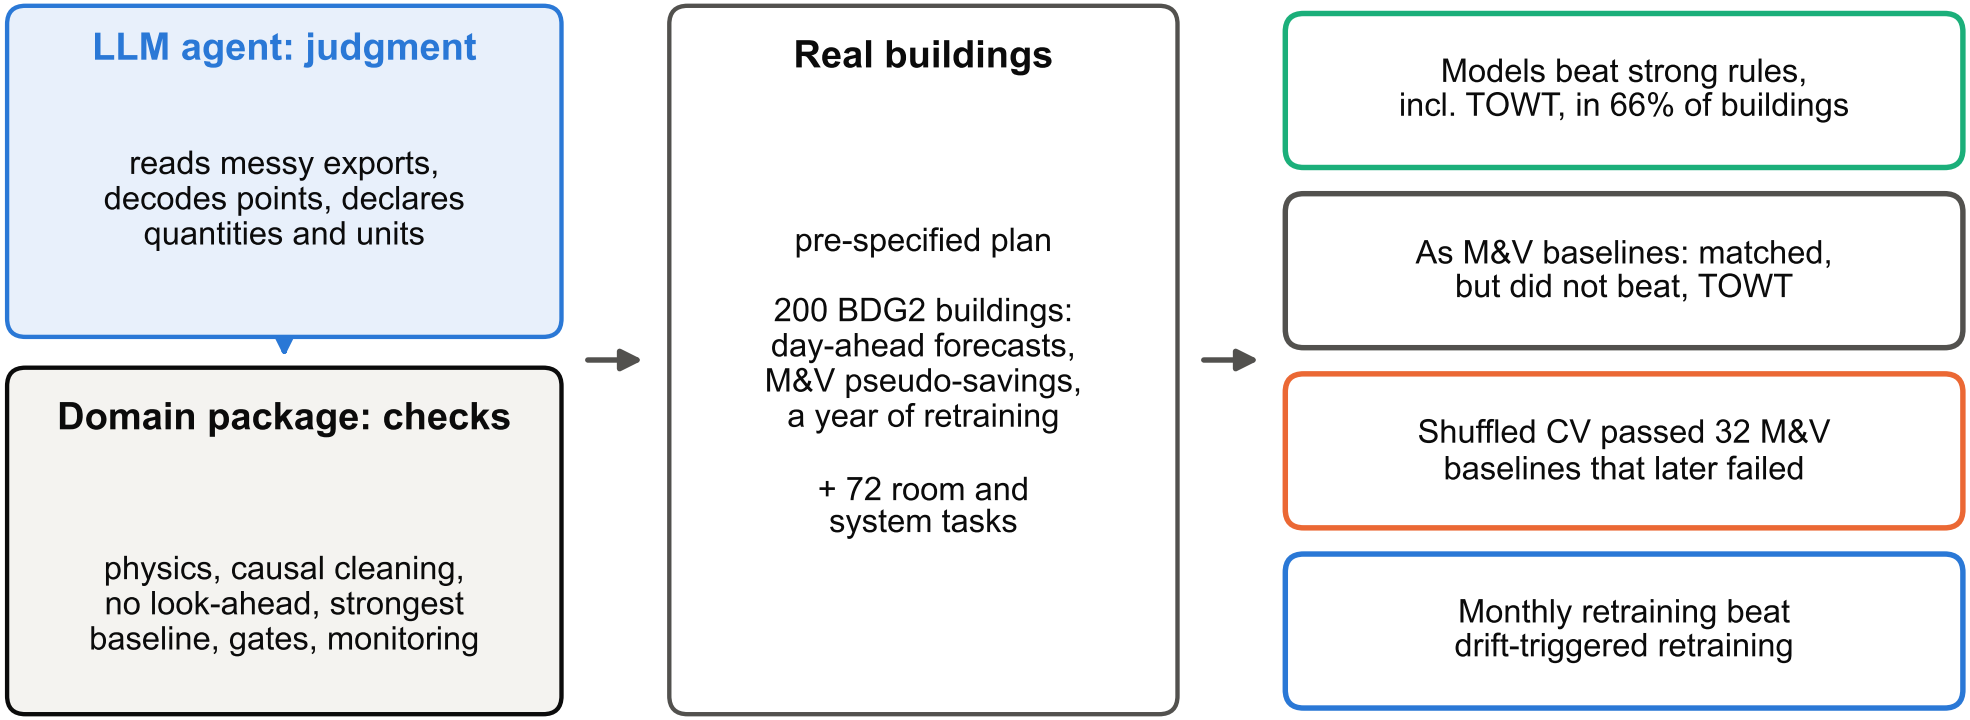
LLM agent: judgment
reads messy exports, decodes points, declares quantities and units
Domain package: checks
physics, causal cleaning, no look-ahead, strongest baseline, gates, monitoring
Real buildings
pre-specified plan
200 BDG2 buildings: day-ahead forecasts, M&V pseudo-savings, a year of retraining
+ 72 room and system tasks
Models beat strong rules, incl. TOWT, in 66% of buildings
As M&V baselines: matched, but did not beat, TOWT
Shuffled CV passed 32 M&V baselines that later failed
Monthly retraining beat drift-triggered retraining